

# PERFORMANCE TUNING WITH SBT-JMH

Scala Hamburg Meetup 20.06.2018

# ABOUT ME

co-founder of inoio

currently doing distributed event sourcing with  
eventuate

on twitter: [@martin\\_grotzke](https://twitter.com/martin_grotzke)

# MICROBENCHMARKING

...done right.

# OPTIMIZATION

...of the right things.

# WHY?

we identified a bottleneck

we want to choose an implementation variant based on  
its performance

just for fun - faster is better!

# MICROBENCHMARKING - THE INTUITIVE APPROACH

```
def bench(name: String, repeat: Int)(fun: => Unit): Unit = {  
    val start = currentTimeMillis()  
    (0 until repeat) foreach (_ => fun)  
    val duration = currentTimeMillis() - start  
    print(s"$name: ${repeat/duration} ops/ms")  
}
```

```
bench("sqrt", 100) {  
    math.sqrt(10000)  
}
```

demo time!

# OH, JIT!

microbenchmarking pitfalls ahead



**INCLUDE JIT ANALYSIS / COMPILATION**

# LOOP OPTIMIZATIONS

# DEAD-CODE ELIMINATION

# CONSTANT FOLDING

# INLINING, DEOPTIMIZATION

# MULTI-THREADING - OH MY!

# MORE PITFALLS

noisy neighbors

gc pauses

classloaders

machine architecture

measuring

# JMH FOR THE RESCUE

Java Microbenchmarking Harness

*JMH is a Java harness for building, running, and analysing nano/micro/milli/macro benchmarks written in Java and other languages targetting the JVM.*

<http://openjdk.java.net/projects/code-tools/jmh/>



# SBT JMH PLUGIN


<https://github.com/ktoso/sbt-jmh>

demo...!

**AWESOME,  
NUMBERS! NOW  
WHAT?!**



# PROFILING?!



source: [survey by zeroturnaround \(2015\)](#)




# ASYNC-PROFILER

*low overhead sampling profiler for Java that does not suffer from Safepoint bias problem. It features HotSpot-specific APIs to collect stack traces and to track memory allocations.*

<https://github.com/jvm-profiling-tools/async-profiler>

# SUPER AWESOME: FLAME GRAPHS!







# USAGE: ASYNC-PROFILER GENERATING FLAME GRAPHS

```
echo 1 > /proc/sys/kernel/perf_event_paranoid  
echo 0 > /proc/sys/kernel/kptr_restrict
```

```
PID=$(jps | grep ForkMain | cut -d" " -f 1)  
profiler.sh -d 30 -f /tmp/flamegraph.svg $PID
```

demo? demo.

# KEY TAKEAWAYS

1. don't build your own `def microbenchmark()`, but use an existing tool!
2. use a low overhead profiler not suffering from safepoint bias.
3. use flame graphs to understand understand your application/code profile.

build - measure - learn

# REFERENCES / FURTHER READING

- Brian Goetz: [Anatomy of a flawed microbenchmark](#)
- Aleksey Shipilëv: [JVM Anatomy Park](#)
- JMH: [openjdk.java.net/projects/code-tools/jmh/](https://openjdk.java.net/projects/code-tools/jmh/)
- mechanical-sympathy: [JMH vs Caliper: reference thread](#)
- sbt-jmh: [github.com/ktoso/sbt-jmh](https://github.com/ktoso/sbt-jmh)
- Nitsan Wakart: [Why \(Most\) Sampling Java Profilers Are Fucking Terrible](#)
- Nitsan Wakart: [The Pros and Cons of AsyncGetCallTrace Profilers](#)
- async-profiler: [github.com/jvm-profiling-tools/async-profiler](https://github.com/jvm-profiling-tools/async-profiler)
- Brendan Gregg: [Flame Graphs](#)

